

# Operating Systems

---

## Chapter 6 Concurrency: Deadlock(死锁) and Starvation(饥饿)

# 6.1 Principles of Deadlock

---

- 6.1 Principles of Deadlock
- 6.2 Deadlock Prevention
- 6.3 Deadlock Avoidance
- 6.4 Deadlock Detection
- 6.5 An Integrated Deadlock Strategy
- 6.6 Dining Philosophers Problem
- 6.7 Summary

# Deadlock

---

- Permanent blocking of a set of processes that either compete for system resources or communicate with each other
- No efficient solution in the general case(通用)
- All deadlock involve conflicting needs for resources by two or more processes (死锁源于两个或者多个进程的资源需求冲突).
  - A common example is the traffic deadlock (交通死锁)

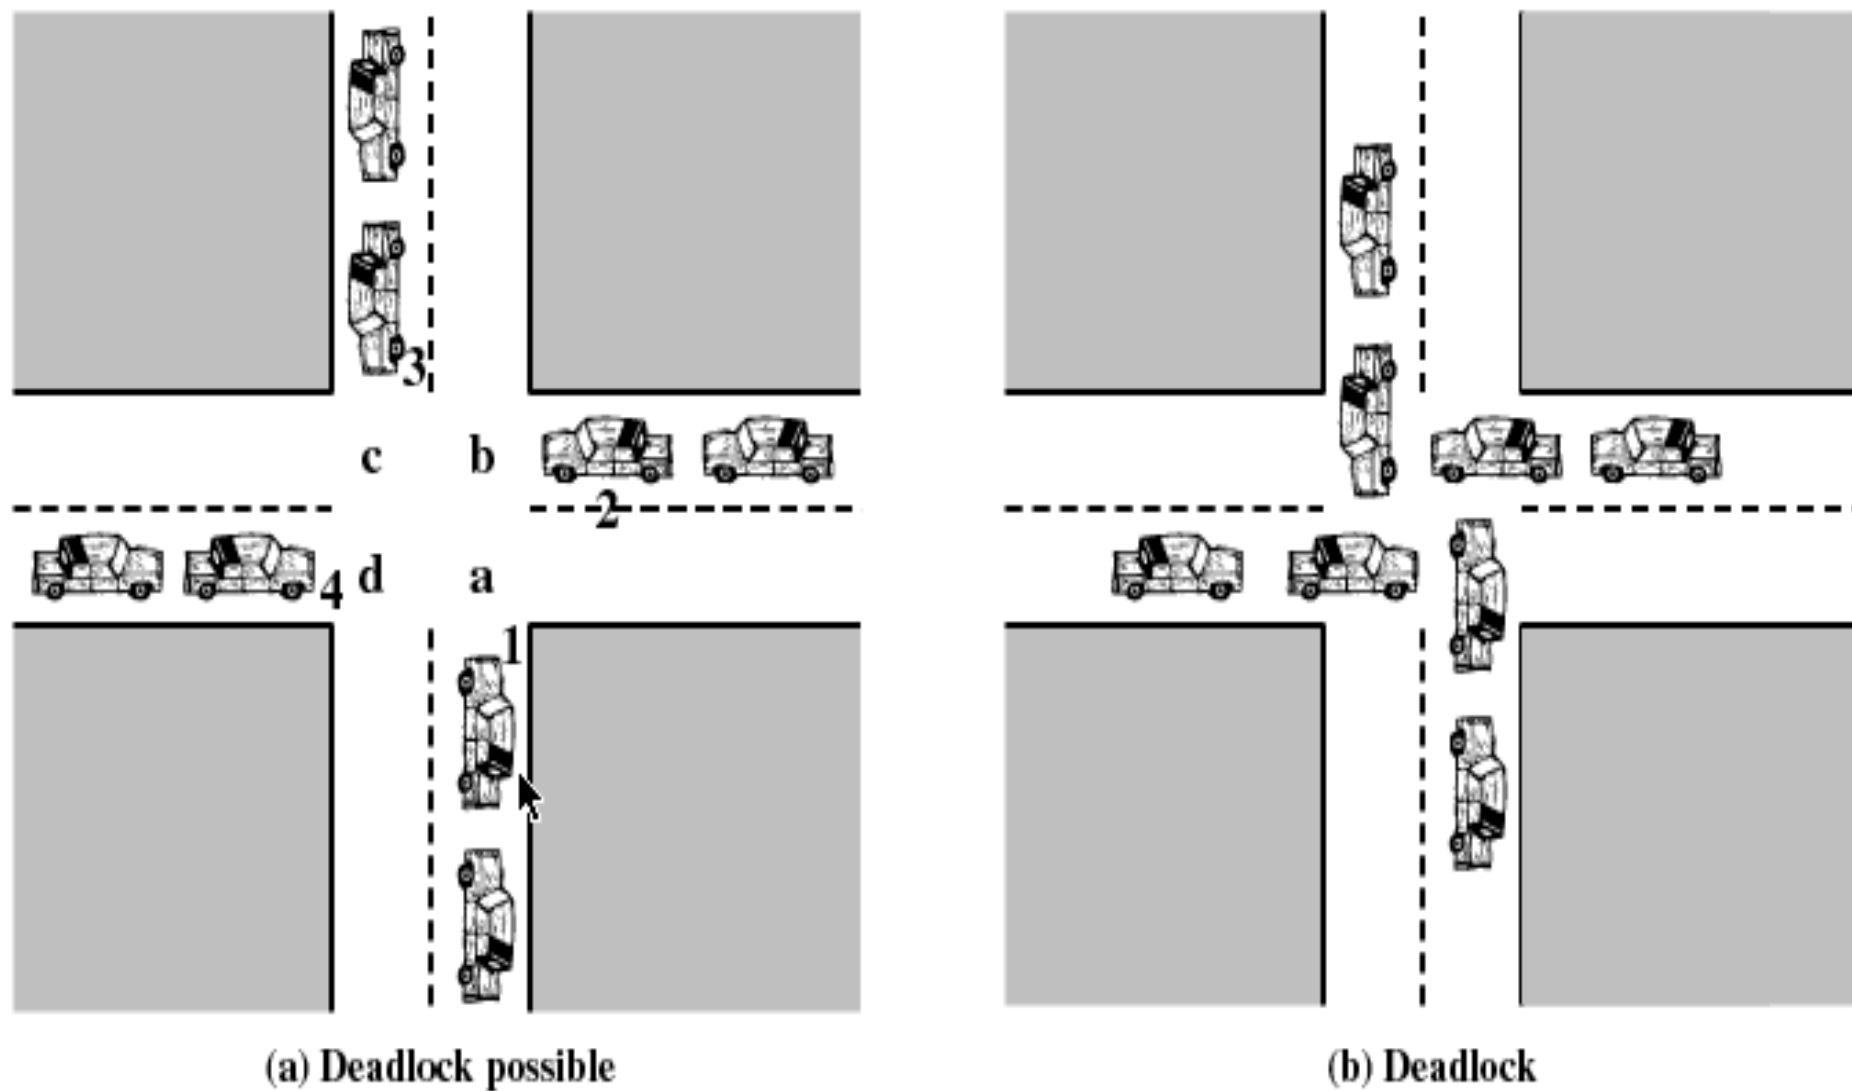




Figure 6.1 Illustration of Deadlock



**Figure 6.2 Example of Deadlock**



**Figure 6.3 Example of No Deadlock [BACO03]**

# Resources Categories(资源的分类)

---

- Reusable Resources(可重用资源)
- Consumable Resources(可消费资源)

# 6.1 Principles of Deadlock

---

- 6.1.1 Reusable Resources
- 6.1.2 Consumable Resources
- 6.1.3 Resource Allocation Graphs
- 6.1.4 The Conditions for Deadlock



# Reusable Resources(可重用资源)

---

- Used by only one process at a time and not depleted(耗尽) by that use
- Processes obtain resources that they later release for reuse by other processes
- Examples include processors, I/O channels, main and secondary memory, devices, and data structures such as files, databases

# Deadlock Example of Reusable Resources

---

Interleaves the execution: p0 p1 q0 q1 p2 q2



Process P		Process Q	
Step	Action	Step	Action
p <sub>0</sub>	Request (D)	q <sub>0</sub>	Request (T)
p <sub>1</sub>	Lock (D)	q <sub>1</sub>	Lock (T)
p <sub>2</sub>	Request (T)	q <sub>2</sub>	Request (D)
p <sub>3</sub>	Lock (T)	q <sub>3</sub>	Lock (D)
p <sub>4</sub>	Perform function	q <sub>4</sub>	Perform function
p <sub>5</sub>	Unlock (D)	q <sub>5</sub>	Unlock (T)
p <sub>6</sub>	Unlock (T)	q <sub>6</sub>	Unlock (D)

**Figure 6.4 Example of Two Processes Competing for Reusable Resources**

# Another Deadlock Example of Reusable Resources

---

- Space is available for allocation of 200Kbytes, and the following sequence of events occur



- Deadlock occurs if both processes progress to their second request

# 6.1 Principles of Deadlock

---

- 6.1.1 Reusable Resources
- 6.1.2 Consumable Resources
- 6.1.3 Resource Allocation Graphs
- 6.1.4 The Conditions for Deadlock

# Consumable Resources(可消费资源)



---

- May be created (produced) and destroyed (consumed) by processes
- Examples include interrupts, signals, messages, and information in I/O buffers

# Deadlock Example of Consumable Resources

---

- Deadlock occurs if receive is blocking



# 6.1 Principles of Deadlock

---

- 6.1.1 Reusable Resources
- 6.1.2 Consumable Resources
- 6.1.3 Resource Allocation Graphs
- 6.1.4 The Conditions for Deadlock

# Resource Allocation Graphs(资源分配图)

- Directed graph(有向图) that depicts(表述) a state of the system of resources and processes



(a) Resource is requested




(b) Resource is held




# Resource Allocation Graphs

---



(c) Circular wait



(d) No deadlock

**Figure 6.5 Examples of Resource Allocation Graphs**

# 6.1 Principles of Deadlock

---

- 6.1.1 Reusable Resources
- 6.1.2 Consumable Resources
- 6.1.3 Resource Allocation Graphs
- 6.1.4 The Conditions for Deadlock


# Conditions for Deadlock(死锁的条件)

---

- Mutual exclusion(互斥)
  - A resource may used only by one process at a time
- Hold-and-wait(占有且等待)
  - A process may hold allocated resources while awaiting assignment of others
- No preemption(非抢占)
  - No resource can be forcibly removed from a process holding it

# Conditions for Deadlock

- Circular wait(循环等待)
  - A closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain



**Figure 6.6** Resource Allocation Graph for Figure 6.1b

# Possibility of Deadlock(死锁的可能性)

---

- Mutual Exclusion
- Hold and wait
- No preemption (抢占权)

# Existence of Deadlock(死锁的存在性)

---

- Mutual Exclusion
- Hold and wait
- No preemption
- Circular wait

# Agenda

---

- 6.1 Principles of Deadlock
- 6.2 Deadlock Prevention
- 6.3 Deadlock Avoidance
- 6.4 Deadlock Detection
- 6.5 An Integrated Deadlock Strategy
- 6.6 Dining Philosophers Problem
- 6.7 Summary

# Deadlock Prevention(死锁预防)

---

- Mutual Exclusion
  - Must be supported by the operating system
- Hold and Wait
  - Require a process request all of its required resources at one time



# Deadlock Prevention

---

- No Preemption
  - Process must release resource and request again
  - Operating system may preempt a process to require it releases its resources
- Circular Wait
  - Define a linear ordering of resource types

# Agenda

---

- 6.1 Principles of Deadlock
- 6.2 Deadlock Prevention
- 6.3 Deadlock Avoidance
- 6.4 Deadlock Detection
- 6.5 An Integrated Deadlock Strategy
- 6.6 Dining Philosophers Problem
- 6.7 Summary

# Deadlock Avoidance(死锁避免)

---

- A decision is made dynamically whether the current resource allocation request will, if granted, potentially lead to a deadlock
- Requires knowledge of future resource requests

# Two Approaches to Deadlock Avoidance

- Process Initiation Denial
  - Do not start a process if its demands might lead to deadlock(如果一个进程的请求会导致死锁，则不启动此进程， *进程启动拒绝*)
- Resource Allocation Denial
  - Do not grant an incremental resource request to a process if this allocation might lead to deadlock(如果一个进程增加资源的请求会导致死锁，则不容许此分配， *资源分配拒绝*)

# Resource Allocation Denial

---

- Referred to as the banker's algorithm
- State of the system(系统状态) is the current allocation of resources to process
- Safe state(安全状态) is where there is at least one sequence that does not result in deadlock
- Unsafe state(不安全状态) is a state that is not safe

# Determination of a Safe State

## a. Initial State

---

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	1	0	0
P2	6	1	2
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	2	2	2
P2	0	0	1
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
0	1	1

Available vector V

(a) Initial state

# Determination of a Safe State

## b.P2 Runs to Completion

	R1	R2	R3
P1	3	2	2
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	1	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	2	2	2
P2	0	0	0
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
6	2	3

Available vector V

(b) P2 runs to completion

# Determination of a Safe State

## c. P1 Runs to Completion

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
7	2	3

Available vector V

(c) P1 runs to completion



# Determination of a Safe State

## d. P3 Runs to Completion

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
9	3	4

Available vector V

(d) P3 runs to completion

# Determination of an Unsafe State

P1 request for one additional unit each of R1 and R3

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix **C**

	R1	R2	R3
P1	1	0	0
P2	5	1	1
P3	2	1	1
P4	0	0	2

Allocation matrix **A**

	R1	R2	R3
P1	2	2	2
P2	1	0	2
P3	1	0	3
P4	4	2	0

**C - A**

R1	R2	R3
9	3	6

Resource vector **R**

R1	R2	R3
1	1	2

Available vector **V**

(a) Initial state

# Determination of an Unsafe State

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix **C**

	R1	R2	R3
P1	2	0	1
P2	5	1	1
P3	2	1	1
P4	0	0	2

Allocation matrix **A**

	R1	R2	R3
P1	1	2	1
P2	1	0	2
P3	1	0	3
P4	4	2	0

**C - A**

R1	R2	R3
9	3	6

Resource vector **R**

R1	R2	R3
0	1	1

Available vector **V**

(b) P1 requests one unit each of R1 and R3

# Deadlock Avoidance Logic(死锁避免逻辑)

```
struct state
{
    int resource[m];
    int available[m];
    int claim[n][m];
    int alloc[n][m];
}
```

(a) global data structures

```
if (alloc [i,*] + request [*] > claim [i,*])
    < error >;                                /* total request > claim*/
else if (request [*] > available [*])
    < suspend process >;
else                                           /* simulate alloc */
{
    < define newstate by:
    alloc [i,*] = alloc [i,*] + request [*];
    available [*] = available [*] - request [*] >;
}
if (safe (newstate))
    < carry out allocation >;
else
{
    < restore original state >;
    < suspend process >;
}
```

(b) resource alloc algorithm

# Deadlock Avoidance Logic

---

```
boolean safe (state S)
{
    int currentavail[m];
    process rest[<number of processes>];
    currentavail = available;
    rest = {all processes};
    possible = true;
    while (possible)
    {
        <find a process  $P_k$  in rest such that
            claim  $[k,*] - \text{alloc } [k,*] \leq \text{currentavail};>$ 
        if (found)                                /* simulate execution of  $P_k$  */
        {
            currentavail = currentavail + alloc  $[k,*]$ ;
            rest = rest -  $\{P_k\}$ ;
        }
        else
            possible = false;
    }
    return (rest == null);
}
```

# Restrictions of Deadlock Avoidance

---

- Maximum resource requirement must be stated in advance
- Processes under consideration must be independent; no synchronization requirements
- There must be a fixed number of resources to allocate
- No process may exit while holding resources

# Agenda

---

- 6.1 Principles of Deadlock
- 6.2 Deadlock Prevention
- 6.3 Deadlock Avoidance
- 6.4 Deadlock Detection
- 6.5 An Integrated Deadlock Strategy
- 6.6 Dining Philosophers Problem
- 6.7 Summary

# Deadlock Detection

---

Reference the textbook for deadlock detection algorithm

	R1	R2	R3	R4	R5
P1	0	1	0	0	1
P2	0	0	1	0	1
P3	0	0	0	0	1
P4	1	0	1	0	1

Request matrix Q

	R1	R2	R3	R4	R5
P1	1	0	1	1	0
P2	1	1	0	0	0
P3	0	0	0	1	0
P4	0	0	0	0	0

Allocation matrix A

R1	R2	R3	R4	R5
2	1	1	2	1

Resource vector

R1	R2	R3	R4	R5
0	0	0	0	1

Available vector

Figure 6.10 Example for Deadlock Detection



# Recovery Strategies Once Deadlock Detected (死锁检测到后的解锁策略)

---

- Abort all deadlocked processes
- Back up (回滚) each deadlocked process to some previously defined checkpoint, and restart all process
  - Original deadlock may occur
- Successively abort (连续取消) deadlocked processes until deadlock no longer exists
- Successively preempt (连续剥夺) resources until deadlock no longer exists

# Selection Criteria Deadlocked Processes (被剥夺或者取消进程的选择标准)

---

- Least amount of processor time consumed so far
- Least number of lines of output produced so far
- Most estimated time remaining
- Least total resources allocated so far
- Lowest priority

# Strengths and Weaknesses of the Strategies

**Table 6.1 Summary of Deadlock Detection, Prevention, and Avoidance Approaches for Operating Systems [ISLO80]**

Approach	Resource Allocation Policy	Different Schemes	Major Advantages	Major Disadvantages
Prevention	Conservative; undercommits resources	Requesting all resources at once	<ul style="list-style-type: none"> <li>• Works well for processes that perform a single burst of activity</li> <li>• No preemption necessary</li> </ul>	<ul style="list-style-type: none"> <li>• Inefficient</li> <li>• Delays process initiation</li> <li>• Future resource requirements must be known by processes</li> </ul>
		Preemption	<ul style="list-style-type: none"> <li>• Convenient when applied to resources whose state can be saved and restored easily</li> </ul>	<ul style="list-style-type: none"> <li>• Preempts more often than necessary</li> </ul>
		Resource ordering	<ul style="list-style-type: none"> <li>• Feasible to enforce via compile-time checks</li> <li>• Needs no run-time computation since problem is solved in system design</li> </ul>	<ul style="list-style-type: none"> <li>• Disallows incremental resource requests</li> </ul>
Avoidance	Midway between that of detection and prevention	Manipulate to find at least one safe path	<ul style="list-style-type: none"> <li>• No preemption necessary</li> </ul>	<ul style="list-style-type: none"> <li>• Future resource requirements must be known by OS</li> <li>• Processes can be blocked for long periods</li> </ul>
Detection	Very liberal; requested resources are granted where possible	Invoke periodically to test for deadlock	<ul style="list-style-type: none"> <li>• Never delays process initiation</li> <li>• Facilitates on-line handling</li> </ul>	<ul style="list-style-type: none"> <li>• Inherent preemption losses</li> </ul>

# Whatever, Deadlock(1/1)

---

- What's it meaning?

- Reasonable?



# Agenda

---

- 6.1 Principles of Deadlock
- 6.2 Deadlock Prevention
- 6.3 Deadlock Avoidance
- 6.4 Deadlock Detection
- 6.5 An Integrated Deadlock Strategy
- 6.6 Dining Philosophers Problem
- 6.7 Summary

# An Integrated Deadlock Strategy(一种综合死锁策略)

---

- Group resources into a number of different resource classes(资源分类)
- Use the linear ordering strategy(线性排序策略) defined previously for the prevention of circular wait to prevent deadlocks between resource classes(类与类之间用线性排序策略避免死锁)
- Within a resource class, use the algorithm that is most appropriate for that class(每个类中使用适合于该类的策略)

# Agenda

---

- 6.1 Principles of Deadlock
- 6.2 Deadlock Prevention
- 6.3 Deadlock Avoidance
- 6.4 Deadlock Detection
- 6.5 An Integrated Deadlock Strategy
- 6.6 Dining Philosophers Problem
- 6.7 Summary

# Dining Philosophers Problem(哲学家就餐问题)




Figure 6.11 Dining Arrangement for Philosophers



# Dining Philosophers Problem (**incorrect**)

---

```
/* program      diningphilosophers */
semaphore fork [5] = {1};
int i;
void philosopher (int i)
{
    while (true)
    {
        think();
        wait (fork[i]);
        wait (fork [(i+1) mod 5]);
        eat();
        signal(fork [(i+1) mod 5]);
        signal(fork[i]);
    }
}
void main()
{
    parbegin (philosopher (0), philosopher (1), philosopher (2),
              philosopher (3), philosopher (4));
}
```

**Figure 6.12**    **A First Solution to the Dining Philosophers Problem**

# Dining Philosophers Problem

---

```
/*program diningphilosophers */
semaphore fork[5] = {1};
semaphore room = {4};
int i;
void philosopher (int I)
{
    while (true)
    {
        think();
        wait (room);
        wait (fork[i]);
        wait (fork [(i+1) mod 5]);
        eat();
        signal (fork [(i+1) mod 5]);
        signal (fork[i]);
        signal (room);
    }
}
void main()
{
    parbegin (philosopher (0), philosopher (1), philosopher (2),
              philosopher (3), philosopher (4));
}
```

**Figure 6.13 A Second Solution to the Dining Philosophers Problem**

# Agenda

---

- 6.1 Principles of Deadlock
- 6.2 Deadlock Prevention
- 6.3 Deadlock Avoidance
- 6.4 Deadlock Detection
- 6.5 An Integrated Deadlock Strategy
- 6.6 Dining Philosophers Problem
- 6.7 Summary